

In [1]:

```
import os
import random
import shutil
import time
import warnings

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim

import torch.utils.data
import torchvision
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import torchvision.models as models
```

In [2]:

```
from torch.cuda.amp import GradScaler
from torch.cuda.amp import autocast
```

In [3]:

```
from torch.utils.tensorboard import SummaryWriter
writer = SummaryWriter()
```

In [4]:

```
!pip install wandb
```

```
Requirement already satisfied: wandb in /opt/conda/lib/python3.8/site-packages (0.12.6)
Requirement already satisfied: shortuuid>=0.5.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (1.0.1)
Requirement already satisfied: sentry-sdk>=1.0.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (1.4.3)
Requirement already satisfied: docker-pycreds>=0.4.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (0.4.0)
Requirement already satisfied: Click!=8.0.0,>=7.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (7.1.2)
Requirement already satisfied: six>=1.13.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (1.16.0)
Requirement already satisfied: GitPython>=1.0.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (3.1.24)
Requirement already satisfied: promise<3,>=2.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (2.3)
Requirement already satisfied: requests<3,>=2.0.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (2.26.0)
Requirement already satisfied: yaspin>=1.0.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (2.1.0)
Requirement already satisfied: psutil>=5.0.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (5.8.0)
Requirement already satisfied: subprocess32>=3.5.3 in /opt/conda/lib/python3.8/site-packages (from wandb) (3.5.4)
Requirement already satisfied: python-dateutil>=2.6.1 in /opt/conda/lib/python3.8/site-packages (from wandb) (2.8.2)
Requirement already satisfied: pathtools in /opt/conda/lib/python3.8/site-packages (from wandb) (0.1.2)
Requirement already satisfied: protobuf>=3.12.0 in /opt/conda/lib/python3.8/site-packages (from wandb) (3.17.3)
Requirement already satisfied: PyYAML in /opt/conda/lib/python3.8/site-packages (from wandb) (5.4.1)
Requirement already satisfied: configparser>=3.8.1 in /opt/conda/lib/python3.8/site-packages (from wandb) (5.0.2)
Requirement already satisfied: gitdb<5,>=4.0.1 in /opt/conda/lib/python3.8/site-packages (from GitPython>=1.0.0->wandb) (4.0.9)
```

Requirement already satisfied: typing-extensions>=3.7.4.3 in /opt/conda/lib/python3.8/site-packages (from GitPython>=1.0.0->wandb) (3.10.0.0)
Requirement already satisfied: smmap<6,>=3.0.1 in /opt/conda/lib/python3.8/site-packages (from gitdb<5,>=4.0.1->GitPython>=1.0.0->wandb) (5.0.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /opt/conda/lib/python3.8/site-packages (from requests<3,>=2.0.0->wandb) (1.26.6)
Requirement already satisfied: charset-normalizer~=2.0.0 in /opt/conda/lib/python3.8/site-packages (from requests<3,>=2.0.0->wandb) (2.0.0)
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.8/site-packages (from requests<3,>=2.0.0->wandb) (2021.5.30)
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.8/site-packages (from requests<3,>=2.0.0->wandb) (3.1)
Requirement already satisfied: termcolor<2.0.0,>=1.1.0 in /opt/conda/lib/python3.8/site-packages (from yaspin>=1.0.0->wandb) (1.1.0)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

In [5]:

```
import wandb
wandb.login()
```

wandb: Currently logged in as: **rochelleli** (use ``wandb login --relogin`` to force relogin)
Out[5]: True

In [6]:

```
wandb.init(project="w251-hw9", entity="rochelleli")
```

Syncing run **frightful-horseman-6** to [Weights & Biases \(docs\)](#).

Out[6]: Display W&B run

In [7]:

```
SEED=2

random.seed(SEED)
torch.manual_seed(SEED)
cudnn.deterministic = True
```

In [8]:

```
torch.cuda.device_count()
```

Out[8]: 1

In [9]:

```
START_EPOCH = 0
```

In [10]:

```
ARCH = 'resnet18'
EPOCHS = 2
```

```
LR = 0.1
MOMENTUM = 0.9
WEIGHT_DECAY = 1e-4
PRINT_FREQ = 10
TRAIN_BATCH=500
VAL_BATCH=500
WORKERS=2
TRAINDIR="/home/ubuntu/data/train"
VALDIR="/home/ubuntu/data/val"
```

```
In [11]: if not torch.cuda.is_available():
        print('GPU not detected.. did you pass through your GPU?')
```

```
In [12]: wandb.init(config={"epochs": EPOCHS, "batch_size": TRAIN_BATCH, "momentum": MOMENTUM, "WEIGHT_DECAY": WEIGHT_DECAY, "arc"
```

Finishing last run (ID:285ki779) before initializing another...

Waiting for W&B process to finish, PID 603... **(success)**.

Synced 4 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Synced **frightful-horseman-6**: <https://wandb.ai/rochelleli/w251-hw9/runs/285ki779>

Find logs at: ./wandb/run-20211031_103311-285ki779/logs

Successfully finished last run (ID:285ki779). Initializing new run:

Syncing run **possessed-crone-1** to [Weights & Biases \(docs\)](#).

Out[12]:

```
In [13]: GPU = 0
```

```
In [14]: torch.cuda.set_device(GPU)
```

```
In [15]: global_step = 0
```

```
In [16]: def train(train_loader, model, criterion, optimizer, epoch):
        global global_step
        batch_time = AverageMeter('Time', ':6.3f')
        data_time = AverageMeter('Data', ':6.3f')
        losses = AverageMeter('Loss', ':.4e')
        top1 = AverageMeter('Acc@1', ':6.2f')
        top5 = AverageMeter('Acc@5', ':6.2f')
```

```

progress = ProgressMeter(
    len(train_loader),
    [batch_time, data_time, losses, top1, top5],
    prefix="Epoch: [{}]" .format(epoch))

# Grad Scaler
scaler = GradScaler()
# switch to train mode
model.train()

end = time.time()
for i, (images, target) in enumerate(train_loader):
    # measure data loading time
    data_time.update(time.time() - end)
    optimizer.zero_grad()

    if GPU is not None:
        images = images.cuda(GPU, non_blocking=True)
    if torch.cuda.is_available():
        target = target.cuda(GPU, non_blocking=True)

    # compute output
    with autocast():
        output = model(images)
        loss = criterion(output, target)

    # measure accuracy and record loss
    acc1, acc5 = accuracy(output, target, topk=(1, 5))
    losses.update(loss.item(), images.size(0))
    top1.update(acc1[0], images.size(0))
    top5.update(acc5[0], images.size(0))

    # compute gradient and do SGD step
    # optimizer.zero_grad()
    # loss.backward()
    # optimizer.step()

    # use the scaler
    scaler.scale(loss).backward()
    scaler.step(optimizer)
    scaler.update()

    # measure elapsed time
    batch_time.update(time.time() - end)
    end = time.time()

    writer.add_scalar("Loss/train", loss, global_step = global_step)
    writer.add_scalar("acc1/train", top1.avg, global_step = global_step)
    writer.add_scalar("acc5/train", top5.avg, global_step = global_step)

wandb.log({"Loss/train": loss, 'acc1/train': top1.avg, 'acc5/train': top5.avg})

```

```
global_step = global_step + 1
```

```
if i % PRINT_FREQ == 0:  
    progress.display(i)
```

In [17]:

```
def validate(val_loader, model, criterion):  
    global global_step  
    batch_time = AverageMeter('Time', ':6.3f')  
    losses = AverageMeter('Loss', ':.4e')  
    top1 = AverageMeter('Acc@1', ':6.2f')  
    top5 = AverageMeter('Acc@5', ':6.2f')  
    progress = ProgressMeter(  
        len(val_loader),  
        [batch_time, losses, top1, top5],  
        prefix='Test: ')
```

switch to evaluate mode
model.eval()

```
with torch.no_grad():  
    end = time.time()  
    for i, (images, target) in enumerate(val_loader):  
        if GPU is not None:  
            images = images.cuda(GPU, non_blocking=True)  
            if torch.cuda.is_available():  
                target = target.cuda(GPU, non_blocking=True)
```

compute output
output = model(images)
loss = criterion(output, target)

measure accuracy and record loss
acc1, acc5 = accuracy(output, target, topk=(1, 5))
losses.update(loss.item(), images.size(0))
top1.update(acc1[0], images.size(0))
top5.update(acc5[0], images.size(0))

measure elapsed time
batch_time.update(time.time() - end)
end = time.time()

```
if i % PRINT_FREQ == 0:  
    progress.display(i)
```

TODO: this should also be done with the ProgressMeter
print(' * Acc@1 {top1.avg:.3f} Acc@5 {top5.avg:.3f}'
 .format(top1=top1, top5=top5))
writer.add_scalar("Loss/val", losses.avg, global_step = global_step)

```

writer.add_scalar("acc1/val", top1.avg, global_step = global_step)
writer.add_scalar("acc5/val", top5.avg, global_step = global_step)

wandb.log({"Loss/val": losses.avg, 'acc1/val': top1.avg, 'acc5/val': top5.avg})

global_step = global_step + 1

return top1.avg

```

In [18]:

```

def save_checkpoint(state, is_best, filename='checkpoint.pth.tar'):
    torch.save(state, filename)
    if is_best:
        shutil.copyfile(filename, 'model_best.pth.tar')

```

In [19]:

```

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self, name, fmt=':f'):
        self.name = name
        self.fmt = fmt
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n
        self.count += n
        self.avg = self.sum / self.count

    def __str__(self):
        fmtstr = '{name} {val' + self.fmt + '} ({avg' + self.fmt + '})'
        return fmtstr.format(**self.__dict__)

```

In [20]:

```

class ProgressMeter(object):
    def __init__(self, num_batches, meters, prefix=""):
        self.batch_fmtstr = self._get_batch_fmtstr(num_batches)
        self.meters = meters
        self.prefix = prefix

    def display(self, batch):
        entries = [self.prefix + self.batch_fmtstr.format(batch)]
        entries += [str(meter) for meter in self.meters]

```

```
print('\t'.join(entries))
```

```
def _get_batch_fmtstr(self, num_batches):  
    num_digits = len(str(num_batches // 1))  
    fmt = '{:' + str(num_digits) + 'd}'  
    return '[' + fmt + '/' + fmt.format(num_batches) + '']'
```

In [21]:

```
def adjust_learning_rate(optimizer, epoch):  
    """Sets the learning rate to the initial LR decayed by 10 every 30 epochs"""  
    lr = LR * (0.1 ** (epoch // 30))  
    for param_group in optimizer.param_groups:  
        param_group['lr'] = lr
```

In [22]:

```
def accuracy(output, target, topk=(1,)):  
    """Computes the accuracy over the k top predictions for the specified values of k"""  
    with torch.no_grad():  
        maxk = max(topk)  
        batch_size = target.size(0)  
  
        _, pred = output.topk(maxk, 1, True, True)  
        pred = pred.t()  
        correct = pred.eq(target.view(1, -1).expand_as(pred))  
  
        res = []  
        for k in topk:  
            correct_k = correct[:k].reshape(-1).float().sum(0, keepdim=True)  
            res.append(correct_k.mul_(100.0 / batch_size))  
        return res
```

In [23]:

```
cudnn.benchmark = True
```

In [24]:

```
import torch.distributed as dist  
import sys
```

In [25]:

```
WORLD_SIZE = 2  
BACKEND = 'nccl'  
  
URL = 'tcp://172.31.17.123:443'  
  
RANK = 1  
  
dist.init_process_group(backend = BACKEND, init_method=URL, rank=RANK, world_size=WORLD_SIZE)
```

```
In [26]: imagenet_mean_RGB = [0.47889522, 0.47227842, 0.43047404]
imagenet_std_RGB = [0.229, 0.224, 0.225]
```

```
In [27]: normalize = transforms.Normalize(mean=imagenet_mean_RGB, std=imagenet_std_RGB)
```

```
In [28]: IMG_SIZE = 224
NUM_CLASSES = 1000
```

```
In [29]: model = models.__dict__[ARCH]()
```

```
In [30]: inf = model.fc.in_features
```

```
In [31]: model.fc = nn.Linear(inf, NUM_CLASSES)
```

```
In [32]: model.cuda(GPU)
model = torch.nn.parallel.DistributedDataParallel(model, device_ids=[GPU])
```

```
In [33]: criterion = nn.CrossEntropyLoss().cuda(GPU)
```

```
In [34]: optimizer = torch.optim.SGD(model.parameters(), LR,
                                     momentum=MOMENTUM,
                                     weight_decay=WEIGHT_DECAY)
```

```
In [35]: scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=EPOCHS)
```

```
In [36]: transform_train = transforms.Compose([
    transforms.Resize((256,256)),
    transforms.RandomCrop(IMG_SIZE, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(imagenet_mean_RGB, imagenet_std_RGB),
])
```

```
In [37]: train_dataset = datasets.ImageFolder(
    TRAINDIR, transform=transform_train)
```



```
In [38]: transform_val = transforms.Compose([
    transforms.Resize((256,256)),
    transforms.RandomCrop(IMG_SIZE, padding=4),
    transforms.ToTensor(),
    transforms.Normalize(imagenet_mean_RGB, imagenet_std_RGB),
])
```

```
In [39]: val_dataset = datasets.ImageFolder(
    VALDIR, transform=transform_val)
```

```
In [40]: train_loader = torch.utils.data.DataLoader(
    train_dataset, batch_size=TRAIN_BATCH, shuffle=False,
    num_workers=WORKERS, pin_memory=True, sampler=torch.utils.data.distributed.DistributedSampler(train_dataset))
```

```
In [41]: val_loader = torch.utils.data.DataLoader(
    val_dataset, batch_size=VAL_BATCH, shuffle=False,
    num_workers=WORKERS, pin_memory=True, sampler=None)
```

```
In [42]: best_acc1 = 0
```

```
In [43]: for epoch in range(START_EPOCH, 2):
    #     adjust_learning_rate(optimizer, epoch)

    # train for one epoch
    train(train_loader, model, criterion, optimizer, epoch)

    # evaluate on validation set
    acc1 = validate(val_loader, model, criterion)

    # remember best acc@1 and save checkpoint
    is_best = acc1 > best_acc1
    best_acc1 = max(acc1, best_acc1)

    save_checkpoint({
        'epoch': epoch + 1,
        'arch': ARCH,
        'state_dict': model.state_dict(),
        'best_acc1': best_acc1,
        'optimizer' : optimizer.state_dict(),
    }, is_best)
```

```

scheduler.step()
print('lr: ' + str(scheduler.get_last_lr()[0]))

writer.add_scalar("lr", scheduler.get_last_lr()[0], global_step = global_step)

wandb.log({'lr': scheduler.get_last_lr()[0]})

```

Epoch: [0][0/1282]	Time 15.616 (15.616)	Data 4.128 (4.128)	Loss 7.0156e+00 (7.0156e+00)	Acc@1 0.00 (
0.00) Acc@5 0.60 (0.60)				
Epoch: [0][10/1282]	Time 0.785 (2.468)	Data 0.127 (0.819)	Loss 6.9182e+00 (6.9759e+00)	Acc@1 0.00 (
0.11) Acc@5 0.80 (0.64)				
Epoch: [0][20/1282]	Time 0.754 (2.112)	Data 0.002 (0.916)	Loss 6.9132e+00 (6.9532e+00)	Acc@1 0.60 (
0.17) Acc@5 1.80 (0.84)				
Epoch: [0][30/1282]	Time 0.764 (1.991)	Data 0.003 (0.945)	Loss 6.8345e+00 (6.9262e+00)	Acc@1 0.00 (
0.23) Acc@5 2.00 (1.00)				
Epoch: [0][40/1282]	Time 0.758 (1.902)	Data 0.002 (0.938)	Loss 6.7156e+00 (6.8819e+00)	Acc@1 0.60 (
0.38) Acc@5 3.20 (1.40)				
Epoch: [0][50/1282]	Time 0.667 (1.863)	Data 0.003 (0.948)	Loss 6.6571e+00 (6.8398e+00)	Acc@1 0.60 (
0.46) Acc@5 2.00 (1.69)				
Epoch: [0][60/1282]	Time 0.695 (1.849)	Data 0.002 (0.968)	Loss 6.5180e+00 (6.7952e+00)	Acc@1 0.60 (
0.52) Acc@5 3.40 (1.95)				
Epoch: [0][70/1282]	Time 0.755 (1.820)	Data 0.089 (0.966)	Loss 6.4503e+00 (6.7562e+00)	Acc@1 1.40 (
0.57) Acc@5 5.20 (2.18)				
Epoch: [0][80/1282]	Time 0.875 (1.809)	Data 0.211 (0.976)	Loss 6.3870e+00 (6.7132e+00)	Acc@1 1.40 (
0.67) Acc@5 4.40 (2.48)				
Epoch: [0][90/1282]	Time 1.269 (1.791)	Data 0.602 (0.976)	Loss 6.3428e+00 (6.6711e+00)	Acc@1 1.40 (
0.77) Acc@5 5.20 (2.80)				
Epoch: [0][100/1282]	Time 1.650 (1.784)	Data 0.982 (0.984)	Loss 6.1171e+00 (6.6282e+00)	Acc@1 2.80 (
0.84) Acc@5 7.00 (3.09)				
Epoch: [0][110/1282]	Time 1.124 (1.769)	Data 0.454 (0.980)	Loss 6.0713e+00 (6.5847e+00)	Acc@1 3.60 (
0.95) Acc@5 9.60 (3.41)				
Epoch: [0][120/1282]	Time 1.230 (1.766)	Data 0.556 (0.987)	Loss 6.0898e+00 (6.5466e+00)	Acc@1 4.00 (
1.04) Acc@5 8.20 (3.70)				
Epoch: [0][130/1282]	Time 1.233 (1.760)	Data 0.557 (0.989)	Loss 5.9793e+00 (6.5097e+00)	Acc@1 1.60 (
1.13) Acc@5 8.40 (3.98)				
Epoch: [0][140/1282]	Time 1.002 (1.750)	Data 0.322 (0.985)	Loss 6.0553e+00 (6.4755e+00)	Acc@1 2.00 (
1.21) Acc@5 8.00 (4.30)				
Epoch: [0][150/1282]	Time 0.981 (1.752)	Data 0.300 (0.993)	Loss 5.9538e+00 (6.4393e+00)	Acc@1 2.80 (
1.30) Acc@5 10.00 (4.59)				
Epoch: [0][160/1282]	Time 0.772 (1.744)	Data 0.091 (0.990)	Loss 5.8895e+00 (6.4059e+00)	Acc@1 2.40 (
1.38) Acc@5 9.80 (4.90)				
Epoch: [0][170/1282]	Time 0.770 (1.738)	Data 0.090 (0.988)	Loss 5.7753e+00 (6.3738e+00)	Acc@1 5.20 (
1.47) Acc@5 9.60 (5.17)				
Epoch: [0][180/1282]	Time 0.773 (1.736)	Data 0.092 (0.988)	Loss 5.8273e+00 (6.3423e+00)	Acc@1 2.80 (
1.58) Acc@5 9.40 (5.45)				
Epoch: [0][190/1282]	Time 0.687 (1.732)	Data 0.002 (0.987)	Loss 5.7511e+00 (6.3109e+00)	Acc@1 3.60 (
1.67) Acc@5 11.40 (5.72)				
Epoch: [0][200/1282]	Time 0.721 (1.735)	Data 0.002 (0.991)	Loss 5.7846e+00 (6.2832e+00)	Acc@1 4.00 (
1.77) Acc@5 11.20 (6.00)				
Epoch: [0][210/1282]	Time 0.690 (1.729)	Data 0.003 (0.987)	Loss 5.6590e+00 (6.2540e+00)	Acc@1 4.20 (

1.87)	Acc@5	11.20 (6.29)										
Epoch: [0][220/1282]	Time	0.689 (1.728)	Data	0.003 (0.987)	Loss	5.6057e+00 (6.2269e+00)	Acc@1	4.60 (1.99)				
1.99)	Acc@5	12.40 (6.56)										
Epoch: [0][230/1282]	Time	0.691 (1.725)	Data	0.002 (0.986)	Loss	5.5602e+00 (6.2001e+00)	Acc@1	6.60 (2.11)				
2.11)	Acc@5	14.20 (6.85)										
Epoch: [0][240/1282]	Time	0.698 (1.722)	Data	0.002 (0.984)	Loss	5.5252e+00 (6.1724e+00)	Acc@1	5.60 (2.22)				
2.22)	Acc@5	14.60 (7.18)										
Epoch: [0][250/1282]	Time	0.694 (1.720)	Data	0.002 (0.983)	Loss	5.5984e+00 (6.1456e+00)	Acc@1	4.40 (2.32)				
2.32)	Acc@5	12.40 (7.47)										
Epoch: [0][260/1282]	Time	0.714 (1.720)	Data	0.002 (0.985)	Loss	5.5894e+00 (6.1197e+00)	Acc@1	4.20 (2.43)				
2.43)	Acc@5	14.60 (7.78)										
Epoch: [0][270/1282]	Time	0.724 (1.717)	Data	0.002 (0.983)	Loss	5.3588e+00 (6.0944e+00)	Acc@1	5.40 (2.54)				
2.54)	Acc@5	16.60 (8.09)										
Epoch: [0][280/1282]	Time	0.714 (1.716)	Data	0.002 (0.983)	Loss	5.3976e+00 (6.0702e+00)	Acc@1	7.00 (2.66)				
2.66)	Acc@5	17.40 (8.37)										
Epoch: [0][290/1282]	Time	0.712 (1.718)	Data	0.003 (0.985)	Loss	5.3315e+00 (6.0474e+00)	Acc@1	7.40 (2.77)				
2.77)	Acc@5	18.00 (8.64)										
Epoch: [0][300/1282]	Time	0.722 (1.714)	Data	0.003 (0.983)	Loss	5.2160e+00 (6.0225e+00)	Acc@1	8.80 (2.91)				
2.91)	Acc@5	20.60 (8.99)										
Epoch: [0][310/1282]	Time	0.734 (1.716)	Data	0.002 (0.985)	Loss	5.2408e+00 (5.9975e+00)	Acc@1	5.60 (3.03)				
3.03)	Acc@5	17.60 (9.30)										
Epoch: [0][320/1282]	Time	0.718 (1.714)	Data	0.003 (0.985)	Loss	5.2504e+00 (5.9747e+00)	Acc@1	7.00 (3.13)				
3.13)	Acc@5	19.60 (9.57)										
Epoch: [0][330/1282]	Time	0.714 (1.713)	Data	0.003 (0.984)	Loss	5.2507e+00 (5.9515e+00)	Acc@1	6.80 (3.25)				
3.25)	Acc@5	19.40 (9.88)										
Epoch: [0][340/1282]	Time	0.683 (1.714)	Data	0.003 (0.986)	Loss	5.0554e+00 (5.9281e+00)	Acc@1	6.80 (3.37)				
3.37)	Acc@5	18.40 (10.16)										
Epoch: [0][350/1282]	Time	0.682 (1.712)	Data	0.002 (0.984)	Loss	5.3180e+00 (5.9071e+00)	Acc@1	8.20 (3.48)				
3.48)	Acc@5	19.60 (10.43)										
Epoch: [0][360/1282]	Time	0.727 (1.713)	Data	0.002 (0.986)	Loss	5.1153e+00 (5.8847e+00)	Acc@1	6.60 (3.60)				
3.60)	Acc@5	20.40 (10.73)										
Epoch: [0][370/1282]	Time	0.694 (1.711)	Data	0.003 (0.985)	Loss	5.1310e+00 (5.8657e+00)	Acc@1	9.40 (3.69)				
3.69)	Acc@5	22.00 (10.97)										
Epoch: [0][380/1282]	Time	0.693 (1.711)	Data	0.002 (0.986)	Loss	4.9869e+00 (5.8452e+00)	Acc@1	8.20 (3.81)				
3.81)	Acc@5	23.80 (11.26)										
Epoch: [0][390/1282]	Time	0.708 (1.711)	Data	0.002 (0.987)	Loss	5.0462e+00 (5.8253e+00)	Acc@1	7.80 (3.93)				
3.93)	Acc@5	21.20 (11.55)										
Epoch: [0][400/1282]	Time	0.713 (1.708)	Data	0.002 (0.984)	Loss	5.0080e+00 (5.8062e+00)	Acc@1	9.20 (4.05)	</			

Epoch: [0][470/1282]	Time 0.690 (1.705)	Data 0.002 (0.985)	Loss 4.8761e+00 (5.6773e+00)	Acc@1 10.60 (
4.83) Acc@5 23.80 (13.68)				
Epoch: [0][480/1282]	Time 0.716 (1.708)	Data 0.002 (0.988)	Loss 4.7287e+00 (5.6586e+00)	Acc@1 13.40 (
4.96) Acc@5 29.60 (13.95)				
Epoch: [0][490/1282]	Time 0.689 (1.706)	Data 0.003 (0.987)	Loss 4.6827e+00 (5.6409e+00)	Acc@1 9.80 (
5.07) Acc@5 28.60 (14.21)				
Epoch: [0][500/1282]	Time 0.683 (1.705)	Data 0.002 (0.986)	Loss 4.8340e+00 (5.6235e+00)	Acc@1 12.60 (
5.19) Acc@5 27.00 (14.47)				
Epoch: [0][510/1282]	Time 0.726 (1.705)	Data 0.003 (0.986)	Loss 4.8993e+00 (5.6076e+00)	Acc@1 10.00 (
5.29) Acc@5 23.80 (14.72)				
Epoch: [0][520/1282]	Time 0.685 (1.704)	Data 0.002 (0.986)	Loss 4.7570e+00 (5.5909e+00)	Acc@1 8.60 (
5.40) Acc@5 25.40 (14.97)				
Epoch: [0][530/1282]	Time 0.691 (1.706)	Data 0.002 (0.988)	Loss 4.6879e+00 (5.5737e+00)	Acc@1 11.40 (
5.51) Acc@5 28.80 (15.23)				
Epoch: [0][540/1282]	Time 0.683 (1.706)	Data 0.002 (0.988)	Loss 4.6027e+00 (5.5574e+00)	Acc@1 13.60 (
5.64) Acc@5 30.40 (15.47)				
Epoch: [0][550/1282]	Time 0.691 (1.705)	Data 0.003 (0.988)	Loss 4.5799e+00 (5.5414e+00)	Acc@1 12.80 (
5.75) Acc@5 31.00 (15.72)				
Epoch: [0][560/1282]	Time 0.684 (1.705)	Data 0.002 (0.988)	Loss 4.5369e+00 (5.5250e+00)	Acc@1 14.80 (
5.86) Acc@5 31.00 (15.98)				
Epoch: [0][570/1282]	Time 0.716 (1.704)	Data 0.002 (0.987)	Loss 4.6132e+00 (5.5088e+00)	Acc@1 11.20 (
5.98) Acc@5 30.80 (16.22)				
Epoch: [0][580/1282]	Time 0.723 (1.706)	Data 0.002 (0.989)	Loss 4.6203e+00 (5.4932e+00)	Acc@1 14.40 (
6.09) Acc@5 30.40 (16.46)				
Epoch: [0][590/1282]	Time 0.687 (1.704)	Data 0.002 (0.988)	Loss 4.8022e+00 (5.4782e+00)	Acc@1 13.00 (
6.21) Acc@5 26.00 (16.68)				
Epoch: [0][600/1282]	Time 0.680 (1.704)	Data 0.002 (0.988)	Loss 4.7113e+00 (5.4627e+00)	Acc@1 12.40 (
6.34) Acc@5 28.60 (16.92)				
Epoch: [0][610/1282]	Time 0.679 (1.703)	Data 0.002 (0.988)	Loss 4.6303e+00 (5.4466e+00)	Acc@1 15.00 (
6.47) Acc@5 29.60 (17.18)				
Epoch: [0][620/1282]	Time 0.781 (1.703)	Data 0.002 (0.987)	Loss 4.4353e+00 (5.4315e+00)	Acc@1 11.40 (
6.57) Acc@5 35.20 (17.41)				
Epoch: [0][630/1282]	Time 0.680 (1.703)	Data 0.002 (0.987)	Loss 4.3729e+00 (5.4163e+00)	Acc@1 16.00 (
6.70) Acc@5 36.20 (17.67)				
Epoch: [0][640/1282]	Time 0.687 (1.702)	Data 0.002 (0.986)	Loss 4.5844e+00 (5.4018e+00)	Acc@1 13.20 (
6.80) Acc@5 32.00 (17.89)				
Epoch: [0][650/1282]	Time 0.682 (1.702)	Data 0.002 (0.987)	Loss 4.3705e+00 (5.3878e+00)	Acc@1 16.60 (
6.91) Acc@5 35.20 (18.14)				
Epoch: [0][660/1282]	Time 0.685 (1.701)	Data 0.002 (0.986)	Loss 4.4043e+00 (5.3734e+00)	Acc@1 14.20 (
7.03) Acc@5 33.60 (18.37)				
Epoch: [0][670/1282]	Time 0.781 (1.701)	Data 0.002 (0.986)	Loss 4.3554e+00 (5.3589e+00)	Acc@1 14.00 (
7.13) Acc@5 33.80 (18.60)				
Epoch: [0][680/1282]	Time 0.685 (1.701)	Data 0.002 (0.986)	Loss 4.3822e+00 (5.3447e+00)	Acc@1 14.60 (
7.24) Acc@5 34.00 (18.84)				
Epoch: [0][690/1282]	Time 0.784 (1.701)	Data 0.002 (0.986)	Loss 4.3757e+00 (5.3306e+00)	Acc@1 16.20 (
7.35) Acc@5 32.00 (19.06)				
Epoch: [0][700/1282]	Time 0.683 (1.701)	Data 0.002 (0.986)	Loss 4.4211e+00 (5.3172e+00)	Acc@1 13.00 (
7.46) Acc@5 31.80 (19.27)				
Epoch: [0][710/1282]	Time 0.686 (1.700)	Data 0.002 (0.985)	Loss 4.5209e+00 (5.3030e+00)	Acc@1 12.00 (
7.57) Acc@5 28.60 (19.51)				
Epoch: [0][720/1282]	Time 0.782 (1.701)	Data 0.002 (0.986)	Loss 4.2508e+00 (5.2898e+00)	Acc@1 16.40 (

[illegible]

Epoch: [0][980/1282]	Time 0.777 (1.695)	Data 0.002 (0.973)	Loss 3.7915e+00 (4.9706e+00)	Acc@1 27.00 (1
0.58) Acc@5 47.00 (25.17)				
Epoch: [0][990/1282]	Time 0.793 (1.695)	Data 0.002 (0.972)	Loss 3.8234e+00 (4.9597e+00)	Acc@1 19.20 (1
0.69) Acc@5 41.80 (25.36)				
Epoch: [0][1000/1282]	Time 0.713 (1.695)	Data 0.002 (0.973)	Loss 3.9626e+00 (4.9495e+00)	Acc@1 19.20 (1
0.78) Acc@5 43.60 (25.54)				
Epoch: [0][1010/1282]	Time 0.779 (1.695)	Data 0.002 (0.973)	Loss 3.8322e+00 (4.9389e+00)	Acc@1 22.60 (1
0.88) Acc@5 46.40 (25.73)				
Epoch: [0][1020/1282]	Time 0.881 (1.694)	Data 0.002 (0.972)	Loss 3.8501e+00 (4.9287e+00)	Acc@1 22.20 (1
0.98) Acc@5 46.80 (25.91)				
Epoch: [0][1030/1282]	Time 0.680 (1.694)	Data 0.003 (0.972)	Loss 3.8828e+00 (4.9180e+00)	Acc@1 22.00 (1
1.09) Acc@5 45.20 (26.10)				
Epoch: [0][1040/1282]	Time 0.880 (1.694)	Data 0.002 (0.971)	Loss 3.7546e+00 (4.9072e+00)	Acc@1 21.00 (1
1.20) Acc@5 46.60 (26.29)				
Epoch: [0][1050/1282]	Time 0.781 (1.694)	Data 0.002 (0.971)	Loss 3.6075e+00 (4.8963e+00)	Acc@1 26.20 (1
1.32) Acc@5 52.20 (26.48)				
Epoch: [0][1060/1282]	Time 0.682 (1.694)	Data 0.002 (0.971)	Loss 3.7612e+00 (4.8858e+00)	Acc@1 21.80 (1
1.42) Acc@5 44.60 (26.67)				
Epoch: [0][1070/1282]	Time 0.698 (1.695)	Data 0.002 (0.972)	Loss 3.7519e+00 (4.8759e+00)	Acc@1 21.60 (1
1.51) Acc@5 45.40 (26.84)				
Epoch: [0][1080/1282]	Time 0.685 (1.695)	Data 0.002 (0.972)	Loss 3.6946e+00 (4.8655e+00)	Acc@1 24.80 (1
1.62) Acc@5 50.80 (27.04)				
Epoch: [0][1090/1282]	Time 0.719 (1.694)	Data 0.002 (0.971)	Loss 3.7519e+00 (4.8554e+00)	Acc@1 22.80 (1
1.72) Acc@5 45.20 (27.21)				
Epoch: [0][1100/1282]	Time 0.717 (1.694)	Data 0.002 (0.971)	Loss 3.8187e+00 (4.8448e+00)	Acc@1 23.60 (1
1.83) Acc@5 46.00 (27.40)				
Epoch: [0][1110/1282]	Time 0.714 (1.694)	Data 0.002 (0.971)	Loss 3.6997e+00 (4.8355e+00)	Acc@1 23.40 (1
1.93) Acc@5 47.40 (27.56)				
Epoch: [0][1120/1282]	Time 0.696 (1.694)	Data 0.002 (0.971)	Loss 3.7813e+00 (4.8253e+00)	Acc@1 21.40 (1
2.04) Acc@5 43.40 (27.74)				
Epoch: [0][1130/1282]	Time 0.779 (1.693)	Data 0.002 (0.971)	Loss 3.7684e+00 (4.8158e+00)	Acc@1 24.60 (1
2.14) Acc@5 45.00 (27.92)				
Epoch: [0][1140/1282]	Time 0.794 (1.693)	Data 0.003 (0.970)	Loss 3.4775e+00 (4.8056e+00)	Acc@1 28.00 (1
2.25) Acc@5 54.00 (28.10)				
Epoch: [0][1150/1282]	Time 0.785 (1.692)	Data 0.002 (0.969)	Loss 3.7545e+00 (4.7956e+00)	Acc@1 23.80 (1
2.35) Acc@5 46.20 (28.28)				
Epoch: [0][1160/1282]	Time 0.782 (1.692)	Data 0.003 (0.969)	Loss 3.8042e+00 (4.7861e+00)	Acc@1 21.80 (1
2.44) Acc@5 45.40 (28.44)				
Epoch: [0][1170/1282]	Time 0.780 (1.692)	Data 0.002 (0.969)	Loss 3.8155e+00 (4.7769e+00)	Acc@1 22.20 (1
2.54) Acc@5 46.40 (28.60)				
Epoch: [0][1180/1282]	Time 0.680 (1.692)	Data 0.002 (0.969)	Loss 3.7569e+00 (4.7673e+00)	Acc@1 24.20 (1
2.65) Acc@5 46.20 (28.77)				
Epoch: [0][1190/1282]	Time 0.692 (1.691)	Data 0.003 (0.969)	Loss 3.7393e+00 (4.7578e+00)	Acc@1 23.60 (1
2.75) Acc@5 50.60 (28.94)				
Epoch: [0][1200/1282]	Time 0.778 (1.690)	Data 0.002 (0.968)	Loss 3.5538e+00 (4.7478e+00)	Acc@1 25.60 (1
2.86) Acc@5 50.20 (29.13)				
Epoch: [0][1210/1282]	Time 0.689 (1.691)	Data 0.002 (0.968)	Loss 3.7215e+00 (4.7378e+00)	Acc@1 21.00 (1
2.97) Acc@5 48.00 (29.31)				
Epoch: [0][1220/1282]	Time 0.688 (1.690)	Data 0.002 (0.968)	Loss 3.6594e+00 (4.7287e+00)	Acc@1 24.40 (1
3.06) Acc@5 47.40 (29.48)				
Epoch: [0][1230/1282]	Time 0.715 (1.690)	Data 0.002 (0.968)	Loss 3.6392e+00 (4.7193e+00)	Acc@1 24.40 (1

```

3.16) Acc@5 47.60 ( 29.64)
Epoch: [0][1240/1282] Time 0.713 ( 1.691) Data 0.002 ( 0.969) Loss 3.5611e+00 (4.7102e+00) Acc@1 25.60 ( 1
3.25) Acc@5 51.40 ( 29.80)
Epoch: [0][1250/1282] Time 0.723 ( 1.690) Data 0.002 ( 0.969) Loss 3.4320e+00 (4.7010e+00) Acc@1 28.20 ( 1
3.35) Acc@5 51.80 ( 29.96)
Epoch: [0][1260/1282] Time 0.721 ( 1.691) Data 0.003 ( 0.969) Loss 3.6723e+00 (4.6922e+00) Acc@1 24.20 ( 1
3.45) Acc@5 47.20 ( 30.11)
Epoch: [0][1270/1282] Time 0.784 ( 1.690) Data 0.003 ( 0.969) Loss 3.5192e+00 (4.6832e+00) Acc@1 25.80 ( 1
3.55) Acc@5 49.80 ( 30.27)
Epoch: [0][1280/1282] Time 0.782 ( 1.690) Data 0.002 ( 0.968) Loss 3.3037e+00 (4.6737e+00) Acc@1 30.20 ( 1
3.67) Acc@5 54.40 ( 30.44)
Test: [ 0/100] Time 7.625 ( 7.625) Loss 3.4868e+00 (3.4868e+00) Acc@1 24.00 ( 24.00) Acc@5 55.00 ( 55.00)
Test: [ 10/100] Time 2.983 ( 2.085) Loss 4.0612e+00 (4.2153e+00) Acc@1 25.00 ( 20.56) Acc@5 43.60 ( 42.16)
Test: [ 20/100] Time 2.837 ( 1.902) Loss 4.0666e+00 (4.0327e+00) Acc@1 10.80 ( 20.34) Acc@5 39.40 ( 43.55)
Test: [ 30/100] Time 2.861 ( 1.831) Loss 4.2806e+00 (4.0831e+00) Acc@1 19.60 ( 19.03) Acc@5 43.00 ( 42.27)
Test: [ 40/100] Time 2.579 ( 1.799) Loss 3.5928e+00 (4.1483e+00) Acc@1 25.20 ( 19.02) Acc@5 51.00 ( 41.46)
Test: [ 50/100] Time 2.898 ( 1.779) Loss 5.1258e+00 (4.1888e+00) Acc@1 9.00 ( 18.22) Acc@5 23.00 ( 40.33)
Test: [ 60/100] Time 2.763 ( 1.767) Loss 3.9949e+00 (4.1866e+00) Acc@1 24.20 ( 18.38) Acc@5 41.00 ( 40.12)
Test: [ 70/100] Time 2.787 ( 1.750) Loss 4.2954e+00 (4.2106e+00) Acc@1 14.00 ( 18.12) Acc@5 36.40 ( 39.71)
Test: [ 80/100] Time 2.922 ( 1.739) Loss 3.9602e+00 (4.2270e+00) Acc@1 22.60 ( 18.09) Acc@5 42.60 ( 39.27)
Test: [ 90/100] Time 2.914 ( 1.729) Loss 4.4120e+00 (4.2401e+00) Acc@1 17.00 ( 17.88) Acc@5 33.20 ( 38.81)
* Acc@1 18.516 Acc@5 39.680
lr: 0.05
Epoch: [1][ 0/1282] Time 6.129 ( 6.129) Data 3.830 ( 3.830) Loss 3.5821e+00 (3.5821e+00) Acc@1 24.00 ( 2
4.00) Acc@5 51.20 ( 51.20)
Epoch: [1][ 10/1282] Time 2.661 ( 2.074) Data 0.003 ( 0.445) Loss 3.5047e+00 (3.4929e+00) Acc@1 26.20 ( 2
5.85) Acc@5 52.60 ( 52.00)
Epoch: [1][ 20/1282] Time 2.507 ( 1.861) Data 0.002 ( 0.238) Loss 3.3250e+00 (3.4390e+00) Acc@1 27.00 ( 2
7.16) Acc@5 55.00 ( 52.63)
Epoch: [1][ 30/1282] Time 2.762 ( 1.803) Data 0.002 ( 0.162) Loss 3.3054e+00 (3.4048e+00) Acc@1 32.60 ( 2
8.21) Acc@5 56.20 ( 53.63)
Epoch: [1][ 40/1282] Time 2.651 ( 1.774) Data 0.002 ( 0.123) Loss 3.3309e+00 (3.3898e+00) Acc@1 29.00 ( 2
8.62) Acc@5 55.40 ( 54.03)
Epoch: [1][ 50/1282] Time 2.240 ( 1.743) Data 0.002 ( 0.100) Loss 3.4286e+00 (3.3757e+00) Acc@1 26.60 ( 2
8.89) Acc@5 53.00 ( 54.12)
Epoch: [1][ 60/1282] Time 1.499 ( 1.724) Data 0.002 ( 0.084) Loss 3.3240e+00 (3.3696e+00) Acc@1 30.40 ( 2
9.13) Acc@5 55.80 ( 54.23)
Epoch: [1][ 70/1282] Time 1.500 ( 1.713) Data 0.002 ( 0.072) Loss 3.2612e+00 (3.3575e+00) Acc@1 28.60 ( 2
9.33) Acc@5 54.80 ( 54.47)
Epoch: [1][ 80/1282] Time 0.778 ( 1.699) Data 0.003 ( 0.064) Loss 3.3153e+00 (3.3532e+00) Acc@1 32.20 ( 2
9.43) Acc@5 59.00 ( 54.62)
Epoch: [1][ 90/1282] Time 0.780 ( 1.696) Data 0.002 ( 0.057) Loss 3.2589e+00 (3.3479e+00) Acc@1 27.60 ( 2
9.39) Acc@5 55.80 ( 54.62)
Epoch: [1][ 100/1282] Time 0.783 ( 1.693) Data 0.002 ( 0.051) Loss 3.0238e+00 (3.3319e+00) Acc@1 35.80 ( 2
9.67) Acc@5 61.00 ( 54.83)
Epoch: [1][ 110/1282] Time 0.683 ( 1.690) Data 0.002 ( 0.047) Loss 3.1109e+00 (3.3257e+00) Acc@1 32.60 ( 2
9.78) Acc@5 60.20 ( 54.99)
Epoch: [1][ 120/1282] Time 0.683 ( 1.684) Data 0.002 ( 0.043) Loss 3.1702e+00 (3.3187e+00) Acc@1 31.80 ( 2
9.88) Acc@5 58.00 ( 55.14)
Epoch: [1][ 130/1282] Time 0.832 ( 1.682) Data 0.092 ( 0.042) Loss 3.1053e+00 (3.3166e+00) Acc@1 31.80 ( 2
9.85) Acc@5 57.60 ( 55.18)

```

Epoch: [1][140/1282]	Time 0.779 (1.682)	Data 0.002 (0.040)	Loss 3.3112e+00 (3.3140e+00)	Acc@1 30.60 (2
9.92) Acc@5 55.40 (55.23)				
Epoch: [1][150/1282]	Time 0.682 (1.682)	Data 0.002 (0.038)	Loss 3.3099e+00 (3.3053e+00)	Acc@1 28.80 (3
0.05) Acc@5 57.00 (55.38)				
Epoch: [1][160/1282]	Time 0.684 (1.683)	Data 0.002 (0.035)	Loss 3.2475e+00 (3.3004e+00)	Acc@1 29.60 (3
0.16) Acc@5 56.40 (55.50)				
Epoch: [1][170/1282]	Time 0.677 (1.682)	Data 0.002 (0.034)	Loss 3.0949e+00 (3.2957e+00)	Acc@1 32.20 (3
0.20) Acc@5 58.80 (55.56)				
Epoch: [1][180/1282]	Time 0.681 (1.680)	Data 0.002 (0.032)	Loss 3.2992e+00 (3.2921e+00)	Acc@1 28.20 (3
0.25) Acc@5 56.80 (55.63)				
Epoch: [1][190/1282]	Time 0.908 (1.678)	Data 0.002 (0.033)	Loss 3.2490e+00 (3.2873e+00)	Acc@1 30.40 (3
0.32) Acc@5 55.40 (55.73)				
Epoch: [1][200/1282]	Time 1.180 (1.677)	Data 0.002 (0.038)	Loss 3.3293e+00 (3.2858e+00)	Acc@1 30.40 (3
0.32) Acc@5 55.00 (55.79)				
Epoch: [1][210/1282]	Time 1.283 (1.676)	Data 0.002 (0.044)	Loss 3.0960e+00 (3.2810e+00)	Acc@1 31.80 (3
0.46) Acc@5 60.60 (55.89)				
Epoch: [1][220/1282]	Time 1.654 (1.675)	Data 0.003 (0.048)	Loss 3.1168e+00 (3.2800e+00)	Acc@1 34.40 (3
0.50) Acc@5 58.40 (55.92)				
Epoch: [1][230/1282]	Time 1.629 (1.675)	Data 0.002 (0.049)	Loss 3.1257e+00 (3.2773e+00)	Acc@1 33.60 (3
0.54) Acc@5 58.20 (55.97)				
Epoch: [1][240/1282]	Time 2.015 (1.676)	Data 0.002 (0.048)	Loss 3.1253e+00 (3.2724e+00)	Acc@1 32.20 (3
0.58) Acc@5 57.60 (56.05)				
Epoch: [1][250/1282]	Time 2.286 (1.677)	Data 0.003 (0.046)	Loss 3.2923e+00 (3.2685e+00)	Acc@1 30.40 (3
0.67) Acc@5 54.00 (56.12)				
Epoch: [1][260/1282]	Time 1.960 (1.677)	Data 0.002 (0.045)	Loss 3.4016e+00 (3.2640e+00)	Acc@1 27.80 (3
0.74) Acc@5 53.80 (56.20)				
Epoch: [1][270/1282]	Time 1.807 (1.676)	Data 0.002 (0.046)	Loss 3.2121e+00 (3.2614e+00)	Acc@1 29.80 (3
0.77) Acc@5 54.80 (56.24)				
Epoch: [1][280/1282]	Time 1.858 (1.675)	Data 0.003 (0.048)	Loss 3.1782e+00 (3.2569e+00)	Acc@1 33.20 (3
0.84) Acc@5 58.60 (56.33)				
Epoch: [1][290/1282]	Time 2.455 (1.676)	Data 0.002 (0.052)	Loss 3.0991e+00 (3.2529e+00)	Acc@1 33.40 (3
0.89) Acc@5 58.20 (56.38)				
Epoch: [1][300/1282]	Time 2.264 (1.675)	Data 0.003 (0.053)	Loss 3.1766e+00 (3.2485e+00)	Acc@1 33.60 (3
0.99) Acc@5 55.40 (56.46)				
Epoch: [1][310/1282]	Time 2.421 (1.675)	Data 0.002 (0.054)	Loss 3.1221e+00 (3.2444e+00)	Acc@1 36.80 (3
1.05) Acc@5 60.20 (56.56)				
Epoch: [1][320/1282]	Time 2.794 (1.676)	Data 0.002 (0.053)	Loss 3.1300e+00 (3.2415e+00)	Acc@1 34.00 (3
1.08) Acc@5 58.40 (56.60)				
Epoch: [1][330/1282]	Time 2.502 (1.675)	Data 0.002 (0.051)	Loss 3.0978e+00 (3.2393e+00)	Acc@1 36.40 (3
1.13) Acc@5 62.00 (56.64)				
Epoch: [1][340/1282]	Time 2.611 (1.675)	Data 0.002 (0.050)	Loss 3.1287e+00 (3.2356e+00)	Acc@1 34.40 (3
1.18) Acc@5 58.00 (56.70)				
Epoch: [1][350/1282]	Time 2.737 (1.675)	Data 0.002 (0.048)	Loss 3.2075e+00 (3.2321e+00)	Acc@1 32.00 (3
1.23) Acc@5 57.60 (56.77)				
Epoch: [1][360/1282]	Time 2.536 (1.674)	Data 0.002 (0.047)	Loss 3.1240e+00 (3.2282e+00)	Acc@1 31.20 (3
1.27) Acc@5 60.00 (56.83)				
Epoch: [1][370/1282]	Time 2.487 (1.674)	Data 0.002 (0.046)	Loss 3.1326e+00 (3.2259e+00)	Acc@1 32.00 (3
1.28) Acc@5 57.20 (56.86)				
Epoch: [1][380/1282]	Time 2.636 (1.673)	Data 0.002 (0.045)	Loss 2.8880e+00 (3.2226e+00)	Acc@1 37.20 (3
1.32) Acc@5 63.00 (56.92)				
Epoch: [1][390/1282]	Time 2.856 (1.673)	Data 0.002 (0.044)	Loss 3.1460e+00 (3.2190e+00)	Acc@1 32.00 (3

[illegible]

Epoch: [1][650/1282]	Time 2.626 (1.676)	Data 0.002 (0.027)	Loss 2.9100e+00 (3.1537e+00)	Acc@1 34.80 (3
2.44) Acc@5 63.20 (58.12)				
Epoch: [1][660/1282]	Time 2.349 (1.676)	Data 0.003 (0.027)	Loss 3.0147e+00 (3.1515e+00)	Acc@1 34.00 (3
2.48) Acc@5 60.40 (58.17)				
Epoch: [1][670/1282]	Time 2.816 (1.676)	Data 0.003 (0.027)	Loss 3.0517e+00 (3.1489e+00)	Acc@1 34.80 (3
2.53) Acc@5 60.60 (58.21)				
Epoch: [1][680/1282]	Time 2.592 (1.676)	Data 0.003 (0.026)	Loss 2.9541e+00 (3.1465e+00)	Acc@1 37.00 (3
2.57) Acc@5 64.40 (58.25)				
Epoch: [1][690/1282]	Time 2.867 (1.676)	Data 0.002 (0.026)	Loss 2.9330e+00 (3.1442e+00)	Acc@1 36.60 (3
2.61) Acc@5 60.60 (58.31)				
Epoch: [1][700/1282]	Time 2.420 (1.676)	Data 0.002 (0.026)	Loss 3.0042e+00 (3.1418e+00)	Acc@1 33.80 (3
2.65) Acc@5 59.80 (58.34)				
Epoch: [1][710/1282]	Time 2.475 (1.675)	Data 0.002 (0.025)	Loss 3.0659e+00 (3.1387e+00)	Acc@1 31.60 (3
2.70) Acc@5 60.20 (58.40)				
Epoch: [1][720/1282]	Time 2.511 (1.675)	Data 0.003 (0.025)	Loss 2.9184e+00 (3.1363e+00)	Acc@1 34.80 (3
2.73) Acc@5 61.80 (58.45)				
Epoch: [1][730/1282]	Time 2.658 (1.675)	Data 0.002 (0.025)	Loss 2.9163e+00 (3.1335e+00)	Acc@1 35.20 (3
2.77) Acc@5 63.20 (58.50)				
Epoch: [1][740/1282]	Time 2.918 (1.675)	Data 0.003 (0.024)	Loss 3.0503e+00 (3.1313e+00)	Acc@1 34.40 (3
2.81) Acc@5 60.40 (58.53)				
Epoch: [1][750/1282]	Time 2.537 (1.676)	Data 0.003 (0.024)	Loss 2.9764e+00 (3.1288e+00)	Acc@1 34.40 (3
2.86) Acc@5 60.20 (58.58)				
Epoch: [1][760/1282]	Time 2.689 (1.676)	Data 0.003 (0.024)	Loss 3.0750e+00 (3.1256e+00)	Acc@1 34.60 (3
2.91) Acc@5 58.60 (58.63)				
Epoch: [1][770/1282]	Time 2.840 (1.676)	Data 0.002 (0.023)	Loss 2.7935e+00 (3.1225e+00)	Acc@1 39.80 (3
2.95) Acc@5 63.20 (58.69)				
Epoch: [1][780/1282]	Time 2.697 (1.675)	Data 0.002 (0.023)	Loss 2.7364e+00 (3.1196e+00)	Acc@1 42.20 (3
3.01) Acc@5 66.20 (58.75)				
Epoch: [1][790/1282]	Time 2.778 (1.675)	Data 0.003 (0.023)	Loss 2.8273e+00 (3.1169e+00)	Acc@1 36.60 (3
3.05) Acc@5 64.40 (58.79)				
Epoch: [1][800/1282]	Time 2.648 (1.675)	Data 0.003 (0.023)	Loss 2.8606e+00 (3.1143e+00)	Acc@1 37.40 (3
3.09) Acc@5 63.60 (58.84)				
Epoch: [1][810/1282]	Time 2.624 (1.676)	Data 0.002 (0.022)	Loss 2.9015e+00 (3.1116e+00)	Acc@1 34.40 (3
3.13) Acc@5 62.40 (58.88)				
Epoch: [1][820/1282]	Time 2.650 (1.675)	Data 0.003 (0.022)	Loss 2.9289e+00 (3.1089e+00)	Acc@1 35.80 (3
3.17) Acc@5 62.00 (58.93)				
Epoch: [1][830/1282]	Time 2.788 (1.675)	Data 0.003 (0.022)	Loss 2.9288e+00 (3.1070e+00)	Acc@1 39.20 (3
3.21) Acc@5 63.20 (58.96)				
Epoch: [1][840/1282]	Time 2.671 (1.676)	Data 0.002 (0.022)	Loss 2.9609e+00 (3.1046e+00)	Acc@1 36.60 (3
3.25) Acc@5 61.40 (59.00)				
Epoch: [1][850/1282]	Time 2.625 (1.675)	Data 0.002 (0.021)	Loss 2.7431e+00 (3.1018e+00)	Acc@1 41.00 (3
3.31) Acc@5 66.40 (59.06)				
Epoch: [1][860/1282]	Time 2.532 (1.675)	Data 0.003 (0.021)	Loss 2.9754e+00 (3.0993e+00)	Acc@1 33.80 (3
3.36) Acc@5 62.20 (59.10)				
Epoch: [1][870/1282]	Time 2.660 (1.675)	Data 0.003 (0.021)	Loss 2.7761e+00 (3.0965e+00)	Acc@1 39.40 (3
3.40) Acc@5 67.00 (59.15)				
Epoch: [1][880/1282]	Time 3.284 (1.675)	Data 0.002 (0.021)	Loss 2.8696e+00 (3.0945e+00)	Acc@1 38.00 (3
3.42) Acc@5 64.40 (59.18)				
Epoch: [1][890/1282]	Time 2.683 (1.676)	Data 0.002 (0.021)	Loss 2.9443e+00 (3.0919e+00)	Acc@1 35.60 (3
3.46) Acc@5 60.00 (59.23)				
Epoch: [1][900/1282]	Time 2.806 (1.676)	Data 0.003 (0.020)	Loss 2.8574e+00 (3.0896e+00)	Acc@1 39.00 (3

3.50)	Acc@5	65.00 (59.28)										
Epoch: [1][910/1282]	Time	2.464 (1.676)	Data	0.002 (0.020)	Loss	2.7528e+00 (3.0872e+00)	Acc@1	40.60 (3				
3.54)	Acc@5	67.60 (59.32)										
Epoch: [1][920/1282]	Time	2.634 (1.676)	Data	0.003 (0.020)	Loss	2.8887e+00 (3.0845e+00)	Acc@1	36.40 (3				
3.58)	Acc@5	60.80 (59.37)										
Epoch: [1][930/1282]	Time	2.769 (1.676)	Data	0.003 (0.020)	Loss	2.8925e+00 (3.0823e+00)	Acc@1	36.40 (3				
3.62)	Acc@5	61.40 (59.41)										
Epoch: [1][940/1282]	Time	2.613 (1.675)	Data	0.002 (0.020)	Loss	2.9197e+00 (3.0801e+00)	Acc@1	36.40 (3				
3.65)	Acc@5	61.00 (59.46)										
Epoch: [1][950/1282]	Time	2.542 (1.675)	Data	0.003 (0.019)	Loss	2.8878e+00 (3.0776e+00)	Acc@1	34.60 (3				
3.69)	Acc@5	61.80 (59.51)										
Epoch: [1][960/1282]	Time	2.648 (1.675)	Data	0.002 (0.019)	Loss	2.9306e+00 (3.0752e+00)	Acc@1	36.00 (3				
3.74)	Acc@5	59.80 (59.55)										
Epoch: [1][970/1282]	Time	2.495 (1.675)	Data	0.002 (0.019)	Loss	2.8924e+00 (3.0728e+00)	Acc@1	34.00 (3				
3.78)	Acc@5	62.00 (59.60)										
Epoch: [1][980/1282]	Time	2.587 (1.675)	Data	0.002 (0.019)	Loss	2.8246e+00 (3.0702e+00)	Acc@1	39.00 (3				
3.83)	Acc@5	64.00 (59.64)										
Epoch: [1][990/1282]	Time	3.048 (1.675)	Data	0.002 (0.019)	Loss	2.7851e+00 (3.0674e+00)	Acc@1	38.00 (3				
3.88)	Acc@5	63.40 (59.70)										
Epoch: [1][1000/1282]	Time	2.512 (1.674)	Data	0.002 (0.019)	Loss	2.9208e+00 (3.0656e+00)	Acc@1	37.40 (3				
3.90)	Acc@5	61.20 (59.73)										
Epoch: [1][1010/1282]	Time	2.742 (1.674)	Data	0.003 (0.018)	Loss	2.8769e+00 (3.0634e+00)	Acc@1	35.80 (3				
3.93)	Acc@5	62.20 (59.77)										
Epoch: [1][1020/1282]	Time	2.478 (1.674)	Data	0.003 (0.018)	Loss	2.8871e+00 (3.0617e+00)	Acc@1	35.80 (3				
3.96)	Acc@5	64.00 (59.80)										
Epoch: [1][1030/1282]	Time	2.512 (1.675)	Data	0.002 (0.018)	Loss	2.7990e+00 (3.0591e+00)	Acc@1	39.00 (3				
4.01)	Acc@5	65.00 (59.84)										
Epoch: [1][1040/1282]	Time	2.765 (1.675)	Data	0.002 (0.018)	Loss	2.7864e+00 (3.0570e+00)	Acc@1	36.00 (3				
4.05)	Acc@5	67.20 (59.88)										
Epoch: [1][1050/1282]	Time	2.521 (1.674)	Data	0.002 (0.018)	Loss	2.7155e+00 (3.0544e+00)	Acc@1	40.40 (3				
4.10)	Acc@5	67.80 (59.93)										
Epoch: [1][1060/1282]	Time	2.503 (1.674)	Data	0.003 (0.018)	Loss	2.7283e+00 (3.0518e+00)	Acc@1	39.00 (3				
4.13)	Acc@5	67.60 (59.99)										
Epoch: [1][1070/1282]	Time	2.816 (1.674)	Data	0.002 (0.018)	Loss	2.8005e+00 (3.0500e+00)	Acc@1	37.80 (3				
4.16)	Acc@5	62.60 (60.01)										
Epoch: [1][1080/1282]	Time	2.694 (1.674)	Data	0.002 (0.017)	Loss	2.7506e+00 (3.0478e+00)	Acc@1	42.00 (3				
4.20)	Acc@5	64.40 (60.05)										
Epoch: [1][1090/1282]	Time	2.512 (1.673)	Data	0.002 (0.017)	Loss	2.7566e+00 (3.0456e+00)	Acc@1	38.40 (3				
4.24)</												

```

Epoch: [1][1160/1282] Time 2.795 ( 1.673) Data 0.002 ( 0.016) Loss 2.9946e+00 (3.0313e+00) Acc@1 36.80 ( 3
4.46) Acc@5 61.40 ( 60.34)
Epoch: [1][1170/1282] Time 2.698 ( 1.673) Data 0.002 ( 0.016) Loss 2.9095e+00 (3.0295e+00) Acc@1 37.80 ( 3
4.50) Acc@5 61.80 ( 60.37)
Epoch: [1][1180/1282] Time 2.105 ( 1.673) Data 0.002 ( 0.017) Loss 2.8175e+00 (3.0272e+00) Acc@1 37.60 ( 3
4.54) Acc@5 63.40 ( 60.42)
Epoch: [1][1190/1282] Time 2.030 ( 1.672) Data 0.002 ( 0.019) Loss 2.8928e+00 (3.0250e+00) Acc@1 40.20 ( 3
4.58) Acc@5 62.40 ( 60.45)
Epoch: [1][1200/1282] Time 2.407 ( 1.672) Data 0.003 ( 0.021) Loss 2.8183e+00 (3.0227e+00) Acc@1 37.60 ( 3
4.62) Acc@5 64.60 ( 60.49)
Epoch: [1][1210/1282] Time 2.407 ( 1.672) Data 0.002 ( 0.022) Loss 2.9153e+00 (3.0200e+00) Acc@1 35.00 ( 3
4.66) Acc@5 63.60 ( 60.54)
Epoch: [1][1220/1282] Time 2.317 ( 1.672) Data 0.002 ( 0.022) Loss 2.9005e+00 (3.0180e+00) Acc@1 37.80 ( 3
4.71) Acc@5 62.20 ( 60.57)
Epoch: [1][1230/1282] Time 2.380 ( 1.673) Data 0.002 ( 0.022) Loss 2.7364e+00 (3.0156e+00) Acc@1 38.80 ( 3
4.74) Acc@5 66.20 ( 60.62)
Epoch: [1][1240/1282] Time 2.232 ( 1.672) Data 0.002 ( 0.021) Loss 2.7297e+00 (3.0136e+00) Acc@1 40.80 ( 3
4.77) Acc@5 66.20 ( 60.65)
Epoch: [1][1250/1282] Time 2.298 ( 1.672) Data 0.002 ( 0.021) Loss 2.5840e+00 (3.0111e+00) Acc@1 41.20 ( 3
4.81) Acc@5 68.20 ( 60.70)
Epoch: [1][1260/1282] Time 2.242 ( 1.672) Data 0.002 ( 0.021) Loss 2.8281e+00 (3.0089e+00) Acc@1 36.40 ( 3
4.85) Acc@5 65.60 ( 60.74)
Epoch: [1][1270/1282] Time 2.237 ( 1.672) Data 0.002 ( 0.021) Loss 2.7000e+00 (3.0069e+00) Acc@1 40.80 ( 3
4.88) Acc@5 64.40 ( 60.78)
Epoch: [1][1280/1282] Time 2.363 ( 1.672) Data 0.002 ( 0.021) Loss 2.5770e+00 (3.0044e+00) Acc@1 41.40 ( 3
4.93) Acc@5 69.00 ( 60.82)
Test: [ 0/100] Time 4.481 ( 4.481) Loss 2.2630e+00 (2.2630e+00) Acc@1 44.20 ( 44.20) Acc@5 76.40 ( 76.40)
Test: [ 10/100] Time 2.920 ( 1.950) Loss 2.3561e+00 (2.8683e+00) Acc@1 49.40 ( 37.80) Acc@5 72.40 ( 63.98)
Test: [ 20/100] Time 2.768 ( 1.820) Loss 2.6376e+00 (2.8284e+00) Acc@1 29.00 ( 35.91) Acc@5 68.40 ( 64.20)
Test: [ 30/100] Time 2.896 ( 1.774) Loss 2.5615e+00 (2.7674e+00) Acc@1 41.40 ( 35.83) Acc@5 67.60 ( 65.27)
Test: [ 40/100] Time 2.577 ( 1.759) Loss 2.7604e+00 (2.7731e+00) Acc@1 39.40 ( 36.46) Acc@5 66.40 ( 65.25)
Test: [ 50/100] Time 2.842 ( 1.749) Loss 3.8991e+00 (2.9221e+00) Acc@1 22.20 ( 34.74) Acc@5 42.40 ( 62.52)
Test: [ 60/100] Time 2.803 ( 1.740) Loss 3.4073e+00 (2.9845e+00) Acc@1 31.60 ( 34.23) Acc@5 54.00 ( 61.42)
Test: [ 70/100] Time 2.888 ( 1.724) Loss 3.5829e+00 (3.0476e+00) Acc@1 26.60 ( 33.56) Acc@5 52.60 ( 60.39)
Test: [ 80/100] Time 2.903 ( 1.717) Loss 3.1231e+00 (3.0971e+00) Acc@1 38.20 ( 33.20) Acc@5 57.20 ( 59.49)
Test: [ 90/100] Time 2.893 ( 1.709) Loss 3.4361e+00 (3.1464e+00) Acc@1 31.80 ( 32.51) Acc@5 53.40 ( 58.62)
* Acc@1 33.104 Acc@5 59.296
lr: 0.0

```

```

In [44]: writer.close()
          %load_ext tensorboard
          %tensorboard --logdir=runs

```

Index of /

Name	Size	Date Modified
.vol/		1/1/20, 12:00:00 AM
Applications/		10/28/21, 12:44:07 AM
bin/		1/1/20, 12:00:00 AM
cores/		6/5/20, 4:55:45 PM
dev/		10/30/21, 11:18:08 AM
etc/		10/15/21, 12:23:15 PM
home/		10/30/21, 11:24:08 PM
Library/		10/15/21, 12:22:52 PM
opt/		2/21/21, 9:52:50 PM
private/		1/1/20, 12:00:00 AM
sbin/		1/1/20, 12:00:00 AM
System/		1/1/20, 12:00:00 AM
tmp/		10/31/21, 4:35:04 AM
Users/		10/15/21, 11:19:49 AM
usr/		1/1/20, 12:00:00 AM
var/		10/15/21, 11:19:51 AM
Volumes/		10/30/21, 11:18:19 AM
.file	0 B	1/1/20, 12:00:00 AM
.VolumeIcon.icns	0 B	

